

Intensive and Extensive Margin Responses in Electricity Conservation: How Households Respond to Financial Rewards

June 6, 2019

Alastair Fraser¹

Preliminary - Please do not circulate or cite

Abstract

Financial rewards are often used to encourage behavioural changes or influence purchasing decisions. In this paper, I investigate how households respond to repeated financial rewards for achieving electricity conservation targets. I track households' extensive margin program participation and intensive margin electricity use across successive conservation challenges offered by a large electrical utility company. Using two complementary empirical strategies, I find that participation in an energy conservation challenge causes households to reduce their electricity use but these reductions rebound close to pre-program levels after program participation ends. I find that households' decisions whether to remain in the program are sensitive to their success or failure in a conservation challenge but, conditional on their success, do not depend on their reductions in electricity use. As a result, households make the opposite decision from what the program design incentivizes; households facing an increased likelihood of achieving their next conservation target are less likely to re-enroll and attempt another conservation challenge. Importantly for designing general incentive programs, this suggests that households use simple heuristics in making decisions rather than incorporating the full amount of information that is readily available to them.

¹The School of Economics at The University of Sydney. I am grateful to Werner Antweiler, Brian Copeland, and Carol McAusland for their instruction and suggestions. Additionally, I thank Patrick Baylis, Joshua Gottlieb, Sumeet Gulati, Thomas Lemeux, Kevin Milligan, Vadim Marmer, Marit Rehavi, Becca Taylor, and participants at a number of workshops and conferences for their assistance. I also thank BC Hydro for providing the electricity data and the UBC Centre for Urban Economics and Real Estate for the housing data used in this paper. Financial support from Environment Canada and Sustainable Prosperity, through the Environmental Economics and Policy Research Network, and the Social Sciences and Humanities Research Council of Canada, is gratefully acknowledged. Email address: alastair.fraser@sydney.edu.au

1 Introduction

Governments and electric utility companies use a wide variety of incentive and information programs to reduce electricity use and improve energy efficiency. The success of these programs has been mixed. Programs can underperform when consumers do not respond to price schedules and information as predicted. Their cost-effectiveness can be reduced when consumers are rewarded for changes that would have occurred in the absence of the program. How consumers use heuristics and deviate from neoclassical models of consumer choice is still poorly understood, and this lack of understanding leads to ad-hoc program designs. Even when energy is conserved or efficiency improved the long-run persistence of changes is often uncertain but important for program effectiveness. In addition, many incentive programs are voluntary. This leaves them vulnerable to selection biases that can make it difficult to evaluate their cost-effectiveness. As a result of these limitations, the widespread deployment of energy efficiency and conservation programs risks generating few and expensive reductions in energy use. These are challenges beyond energy use; ensuring that programs cost-effectively deliver anticipated results requires improved models of how incentives and information affect decisions and the careful evaluation of existing programs

In this paper, I analyze a novel program that incentivizes energy conservation through repeated financial rewards. BC Hydro's Team Power Smart program offers households the opportunity to undertake annual electricity conservation "challenges." Households were told that if they successfully reduce their annual electricity use by 10% relative to their use in the previous 12 months they would receive a \$75 financial reward. This reward is worth 10% of the average household's annual electricity bill. After each annual conservation challenge, all households have the option of participating in an additional conservation challenge. Importantly, subsequent conservation challenges require an additional 10% reduction in electricity use relative to the previous year and are available to households regardless of their past challenge success or failure. This allows me to track households' extensive margin program participation and intensive margin electricity use decisions across multiple conservation challenges. This provides a rich setting for evaluating how consumers respond to repeated financial rewards and information on their success or failure.

I make two important contributions. First, using an event study empirical strategy I estimate the short-run reductions and long-run persistence of changes in electricity use

associated with participation in Team Power Smart. I find that an initial conservation challenge is associated with an immediate 4.4% average reduction in electricity use which lasts throughout the twelve months of the initial challenge. In addition, electricity use continues to decline among households that re-enroll and rebounds close to pre-program levels as households end their participation in the program, and this paper finds minimal evidence of heterogeneous effects across household characteristics.

The rebound in electricity use and the lack of heterogeneity contrasts with the findings of [Ito \(2015\)](#), who evaluates financial rewards for energy conservation offered for four months under the California 20/20 program and finds both persistence and heterogeneity in treatment effects.² The findings in this paper suggest that whether households respond to conservation targets with persistent reductions may depend on the way electricity is used within a home and vary by region. Alternatively, the context in which a financial reward program is offered—a mandatory program during an electricity crisis as in [Ito \(2015\)](#), or a voluntary program as part of routine electricity use as in this paper—may be important to whether households respond with short-run reductions or persistent changes. This paper’s finding of continued reductions in electricity use when households re-enroll, and a rebound when they don’t, is similar to the “action and backsliding” found by [Allcott and Rogers \(2014\)](#) in households’ responses to repeated home energy reports.

Despite this conservation program being voluntary, I find no evidence that households are responding strategically in any way that would contribute to the well known ‘additionality’ problem ([Boomhower and Davis, 2014](#)). In contrast, self-selection can increase a voluntary programs cost-effectiveness by raising the share of participant households that respond to the incentive. This is particularly important for electricity conservation as household electricity use, even after controlling for weather changes, exhibits large idiosyncratic variations over time. As a result, even households that do not respond to the incentive may still receive rewards for sufficiently large idiosyncratic reductions in electricity use, and this share can be substantial in mandatory programs.

²Several other papers have considered financial rewards and energy conservation. [Dolan and Metcalfe \(2015\)](#) undertake a randomized controlled trial (RCT) and find large financial rewards cause a large conservation over the two months of their treatment period. They find persistent effects over the two post-program months they observe. Several papers, primarily from the psychology literature, have also undertaken randomized control trials of financial rewards for energy conservation ([Mizobuchi and Takeuchi, 2012](#); [Midden et al., 1983](#); [McClelland and Cook, 1980](#); [Winett et al., 1978](#)). However, these consider very short timeframes, small sample sizes, and unrepresentative electricity users and are of limited use in understanding how households respond to financial rewards. [Gerard and Costa \(2015\)](#) study a suite of mandatory incentives introduced in response to an electricity supply crisis in Brazil, including financial rewards, but cannot disentangle the effect of different policies.

Second, I find that households’ decisions on whether to remain in the program are sensitive to their success or failure in a conservation challenge, but are insensitive to the detailed information on their actual conservation effort that is provided to them.³ This creates a discontinuity in the probability they continue to an additional challenge, and as a result households make the opposite decision from what the program design incentivizes; households facing an increased likelihood of achieving their next conservation target are less likely to re-enroll and attempt another conservation challenge. Importantly, the program defined success using a threshold of a 9.5% reduction while households were expecting this to occur at 10%. This allows me to demonstrate that the fuzzy regression discontinuity approach remains valid despite bunching by households at the 10% reduction target. Using this empirical strategy I find that an additional conservation challenge causes lower electricity use. This is consistent with the event study results and shows that voluntary financial reward programs can cause reductions in electricity use, but that electricity use tends to rebound in the absence of the financial incentive.

Information provision is a ubiquitous feature of energy conservation and other incentive programs, both as a stand alone intervention and bundled with incentives such as price changes or financial rewards. This paper’s findings suggest that while households are responding to information, they either do not take the full set of information into account or do not respond to it as a standard neoclassical model of well informed consumer decisions would predict.⁴ This suggests that households are using a simple heuristic in making their participation decisions rather than incorporate the detailed information that is readily available to them.

The use of heuristics and inattention to information, whether rational or a sub-optimal response to the incentive, may be a pervasive feature of consumers’ responses and is an important consideration in the design of incentive programs. This is important for programs based on providing consumers with detailed information. For example, [Schleich et al. \(2017\)](#) undertakes a RCT where households can access detailed information on their

³This information is provided to participants both through a letter at the end of their conservation challenge and through an online portal. The majority of households regularly log into their online portal ([Kassirer et al., 2014](#)) so their lack of response is not due to a lack of awareness of their progress or success.

⁴An important model of decision making in Psychology is self-efficacy, which refers to an individual’s beliefs about their ability to affect outcomes through their actions ([Bandura, 1977](#)). Households could observe their success, interpret this as evidence they can affect the outcome of electricity conservation, and decide to re-enroll due to their updated beliefs on their self-efficacy. However, the sharp discontinuity in this setting—and the information provided to households on their actual conservation achieved—precludes the concept of self-efficacy from alone explaining households’ behaviour.

consumption through an online portal or through mailed reports and finds this is effective in persistently reducing electricity use. [Jessoe and Rapson \(2014\)](#) use a RCT to find in-home displays on the price and quantity of electricity used increases consumers price-elasticity, while [Martin and Rivers \(2016\)](#) show how similar in-home displays can directly cause energy conservation.⁵ This paper’s findings are also consistent with previous work that suggests the arrival of information, not only the information content itself, can serve as a nudge or reminder on electricity use. [Gilbert and Graff Zivin \(2013\)](#) show that simply receiving an electricity bill can reduce electricity use in the short term by 1%, and [Sexton \(2015\)](#) finds that increased billing frequency lowers electricity use. The responsiveness of households to their passing or failing a challenge but not the incentive structure is also suggestive that intrinsic motivations for conserving energy may be important ([Gneezy et al., 2011](#)). Previous work considering intrinsic motivations include the home energy reports studied in [Allcott \(2011\)](#) and [Allcott and Rogers \(2014\)](#), and [Ito et al. \(2015\)](#) who directly compare moral suasion against the incentive of short-term higher prices through two RCTs. They find that while both interventions cause lower short-run electricity use, only the short-term price increase generated persistent effects. Information also plays an important role in programs that change a price schedule or offer a financial reward. These programs assume consumers are sufficiently informed of their own energy use and the program design to respond to the incentive as intended; programs can fail to deliver anticipated outcomes if this is not the case. For example, [Ito \(2014\)](#) exploits a spatial discontinuity among households experiencing different block pricing schedules and finds that consumers respond to average prices instead of marginal prices. He concludes that the nonlinear price schedule is unsuccessful in achieving energy conservation due to this suboptimal response.

The remainder of this paper is organized as follows. Section 2 describes the institutional setting, design of the Team Power Smart electricity conservation program, and data. Section 3 gives an overview of the two empirical approaches used in this paper. Section 4 describes the event study empirical strategy and results. Section 5 explores households’ decisions whether to leave the program, describes the instrumental variables strategy, and presents its results. Section 6 discusses the cost effectiveness of the program, and I conclude in Section 7.

⁵There is no a-priori reason that additional information would lead to decreases rather than increases in consumption; [Wichman \(2017\)](#) find that switching from bi-monthly to monthly billing for water increased its use.

2 Institutional Setting, Program Design, and Data

BC Hydro is Canada’s second largest integrated electrical utility company and serves 1.7 million residential customers covering 95% of the population in British Columbia (BCH, 2014). BC Hydro is owned and has a mandate set by the provincial government of British Columbia. The B.C. government, through its Clean Energy Act, has required that BC Hydro “[meet] at least 66 per cent of the expected increase in demand through conservation and efficiency by 2020” (BCH, 2014). As part of their efforts to achieve this mandate and in response to the Conservation Potential Review, BC Hydro launched a new conservation initiative—a program targeting ongoing behaviours called Team Power Smart.⁶ Team Power Smart is a voluntary program promoted and summed up by BC Hydro with “*Looking to save money on your electricity bills? Become a member of Team Power Smart and challenge yourself to reduce your home’s electricity use by 10% in the next year. If you’re successful, you can earn a [\$75] reward.*”⁷ This electricity conservation challenge requires households to reduce their aggregate electricity use over the 12-month challenge by 10% relative to a conservation target. Each household’s conservation target is their own annual electricity use over the preceding 12 months, adjusted for changes in heating degree days to help prevent households from being unduly penalized or rewarded for changes in weather. Beginning a challenge requires a minimal time cost of registering online. Households can start a challenge in any month of the year as long as they have 12 months of electricity use in their current home to establish their target. Online signup ensures that all participants can view their progress towards their conservation target through the BC Hydro website and access a variety of tips and suggestions for reducing their electricity use. The online account provides households with feedback on both their monthly and cumulative progress towards their annual 10% conservation target.⁸ Because it is the aggregate annual conservation that matters for success in a challenge, households can miss their 10% target in any month and still pass the challenge.

Upon completing the 12 months of the challenge BC Hydro undertakes a final evaluation of household’s cumulative conservation. BC Hydro applies the final weather adjustment,

⁶It is possible for households to join Team Power Smart to view their electricity use online without undertaking a conservation challenge. For simplicity, I will use Team Power Smart to refer to those households which also undertake a conservation challenge. I do not observe households which registered online without undertaking a conservation challenge.

⁷BC Hydro Team Power Smart website landing page. The reward value for challenges studied in this paper is \$75 which was reduced to \$50 in September 2014. I exclude households undertaking a challenge under the new \$50 reward value. Accessed June 2017.

⁸An example of the online portal is shown in Appendix Figures A.1a and A.2b.

accounts for bi-monthly billing and any idiosyncratic factors, and evaluates whether the household passed or failed their challenge. While the conservation target advertised to customers is 10%, BC Hydro evaluates final success or failure against a 9.5% conservation threshold. Households that reduced their electricity use by greater than or equal to 9.5% below their target pass their challenge while the rest fail. BC Hydro notifies all customers of how much their electricity use changed alongside their success or failure and gives successful households the choice of a rebate through either a cheque or credit applied to their account.

Structure of Additional Conservation Challenges

A novel feature of Team Power Smart is that all households have the option of re-enrolling in additional annual conservation challenges. Upon completion of each challenge both households that pass and fail their challenge are given the same option to start a subsequent conservation challenge for another \$75 rebate. Each subsequent challenge follows the same process as the initial conservation challenge; households have a goal of another 10% conservation target measured and weather adjusted relative to their previous 12 months of electricity use. The new reduction target is independent of whether the prior 12 months contained a challenge or not, and independent of whether the prior challenge was a success or failure. The baseline for a household immediately starting an additional challenge would be the 12 months of the just completed challenge, while a household waiting 4 months before starting their next challenge would have a baseline set by the average of their last 8 months of their previous challenge and the 4 months of the gap prior to starting their next challenge. Because each additional 10% conservation challenge is evaluated relative to the prior 12 months, the reduction in electricity use achieved by a household during a challenge affects their incentives on when and whether to undertake a subsequent challenge. Under the reasonable assumption of increasing marginal costs to electricity conservation, the greater the conservation achieved during a challenge the greater the incentive to postpone a subsequent challenge or leave the program.

Data and Household Characteristics

Under a non-disclosure agreement with BC Hydro I obtained an anonymized sample of monthly electricity billing records for 10,000 Team Power Smart program participants

and 20,000 non-participants from January 2006 to December 2015.⁹ The panel includes customers' Team Power Smart program participation history including the number of conservation challenges, each challenges start and end date, whether the challenge was successful, and the building and heating type of the household. Individual household characteristics from BC Assessment including building type, number of bedrooms, assessed value, floor space, and the postal codes Forward Sortation Area were merged with the BC Hydro panel. Removing duplicate accounts, erroneous data, dropping households with electricity use more than 5 standard deviations from the mean, and removing households that joined Team Power Smart after the reward value was changed in September 2014, left a sample of 8,877 households participating in Team Power Smart.

Average monthly electricity bills among participant households are \$62—making the rebate reward of \$75 equivalent to 10% of a household's annual electricity bill, in addition to their bill savings. Differences in electricity use between participants and non-participants is almost entirely a composition effect; after controlling for building and heating type the average electricity use among participant households is 9.3 kWh (p-value 0.08), or 1.1%, higher than non-participants during the pre-program year of 2006.¹⁰

Outcomes During Multiple Conservation Challenges

Table 1 summarizes the decisions and outcomes of participant households across multiple conservation challenges. During the initial three challenges approximately 61% of households decide to re-enroll in an additional challenge.¹¹ Consistent with an increasing difficulty of achieving additional reductions in electricity use, the unconditional probability of passing a conservation challenge declines with additional challenges.¹² Households are more likely to re-enroll in another challenge if they pass, rather than fail, their current challenge. In contrast, households are less likely to pass their next challenge if they passed

⁹I define participant households as those which participate in Team Power Smart prior to the panel end in December 2015, and non-participant households as those which do not participate. The sample was selected from households which did not move over the panel period.

¹⁰The idiosyncratic process of advertising the voluntary Team Paper Smart program means differences in household characteristics between participant and non-participant households are not necessarily due to self-selection. These differences are described in Appendix A.2.

¹¹The probability of re-enrolling declines in part mechanically among later challenges due to the limited panel length.

¹²This decline is statistically significant. Regressing an indicator for whether households passes their challenge on the Challenge number finds that decline in the probability of passing is statistically significant at the 1% level, (coefficient on *Challenge* of -0.0212, t-statistic of -9.91), but loses significance if restricted to challenges 6 through 9 (coefficient on *Challenge* of 0.0067, t-statistics 0.28).

Table 1: **Probability of Challenge Outcomes**

Challenge number	HH's in Chal.	Probability Of Re-Enrolling			Probability Of Passing		
		All	if Failed Challenge	if Passed Challenge	All	if Failed Prev. Chal.	if Passed Prev. Chal.
1	8,877	0.62	0.55	0.77	0.34		
2	5,531	0.60	0.56	0.71	0.31	0.33	0.29
3	3,346	0.60	0.56	0.71	0.28	0.30	0.24
4	2,014	0.54	0.51	0.64	0.26	0.28	0.23
5	1,091	0.46	0.41	0.60	0.24	0.27	0.17
6	498	0.38	0.36	0.44	0.24	0.25	0.21
7	188	0.27	0.26	0.28	0.29	0.28	0.31
8	50	0.12	0.07	0.33	0.18	0.20	0.13
9	6	0.00	0.00	0.00	0.33	0.33	0.33

Notes: HH's in Chal. is the number of households undertaking their first, second, etc, challenge. Probability of Re-Enrolling is the probability of re-enrolling in a subsequent conservation challenge, conditional on being in the current challenge. Probability of Re-Enrolling if Failed [Passed] Challenge is the probability of re-enrolling conditional on failing [passing] the current challenge. The Probability of Passing is for a household's current challenge, while the Probability of Passing if Failed [Passed] Prev. Chal. is the probability of passing the current challenge conditional on the Fail or Pass status of the previous challenge.

their previous challenge.¹³ This pattern matches the incentive structure previously discussed; passing a challenge requires achieving the 9.5% conservation target, which makes passing the next challenge harder compared to not achieving the initial target.¹⁴

3 Empirical Strategies

There are three principal challenges to estimating the causal effect of repeated conservation challenges. First, as in all voluntary programs, households may self-select into Team Power Smart based on observable and unobservable time invariant characteristics. This self-selection could make non-participant households an unsuitable counterfactual for electricity use among participant households, had they not participated in the pro-

¹³This difference, from Table 1, between the probability of passing conditional on a households success and the probability conditional on failure in the previous challenge is statistically significant. Regressing an indicator for passing the challenge on an indicator for success in the previous challenge finds that prior success is associated with a 4.65% (t-statistic -5.69) lower average probability of success in the following challenge. This difference is of a similar magnitude and statistically significant at the 5% level for individual Challenges 2-5, after which differences lose statistical significance.

¹⁴As households choose the start date of subsequent conservation challenges they could strategically establish a new baseline before undertaking their next challenge, or sign up after periods of unusual weather. In Appendix A.3 I find no evidence this is occurring.

gram. Secondly, households may start their first conservation challenge based on shocks to their past electricity consumption or expectations of their future consumption. For example, households may select into the program in response to a particularly cold winter which caused a large electricity bill. Households may also take advantage of anticipated reductions in their electricity use such as the purchase of an efficient dryer or leaving on holiday. By signing up in advance or in conjunction with their anticipated reduction in electricity use, a household could receive credit for reductions in electricity use that were not caused by the conservation program. Lastly, all households have the option of continuing to additional conservation challenges. This makes the persistence of energy savings and the causal effect of subsequent conservation challenges dependent on the households' decisions to select into additional challenges.

To address these challenges I employ two complementary empirical strategies. In Section 4 I use an event study research design to estimate the monthly average changes in electricity use associated with participation in Team Power Smart. By estimating changes over time within households, this strategy identifies program effects independent of self-selection into the program on observable and unobservable time invariant characteristics. For clarity, I use “program effects” to refer to changes in electricity use that are associated with participation in conservation challenges but which may or may not be causally due to the program. I restrict “treatment effect” to refer to standard causal effects. Plotting these estimates provides a visual time trend of the changes in electricity use leading up to the initial conservation challenge, during the months of each challenge, and over the months after a household leaves the program. These trends provide insight into potential self-selection into the program based on households' expectations of future electricity use and past shocks to their consumption, along with the persistence of reductions in electricity use and the effect of subsequent conservation challenges. This strategy provides detailed information on households' electricity conservation decisions and self-selection into the program, and finds no evidence that reductions in electricity associated with participation are not causally due to the program. However, because self-selection cannot be ruled out it cannot strictly identify causal treatment effects.

To identify causal treatment effects, Section 5 exploits a discontinuity in the probability that households continue to a second conservation challenge. This discontinuity is caused by households that just succeed in their conservation challenge being significantly more likely to continue to a second challenge compared to those households that just failed their challenge. I use this discontinuous probability change in a fuzzy Regression Discon-

tinuity design which, under several identifying assumptions, identifies the Local Average Treatment Effect (LATE) of a second conservation challenge despite self-selection on time-invariant characteristics, expectations of future electricity use, or shocks to households' past consumption (Lee and Lemieux, 2010). By comparing households' decisions whether to re-enroll in a subsequent challenge to their intensive margin reductions in electricity use, Section 5 also provides insights into how households respond to information on their conservation challenge success.

4 Event Study Empirical Strategy

The general event study model is of the form

$$y_{it} = \sum_{\tau=-T}^T \beta_{\tau} D_{i,t-\tau+1} + \alpha_i + d_t + \epsilon_{it} \quad (1)$$

where y_{it} is the log of monthly electricity use for individual i in month t , α_i is an individual fixed effect, and d_t is an indicator for date t . $D_{i,t-\tau+1}$ is a dummy variable equaling 1 if individual i in month t began treatment in month $t-\tau+1$, where τ is the measure of event time in months.¹⁵ I define the start of treatment as the month a household undertakes its first conservation challenge, $\tau = 1$. β_{τ} are the non-parametric program effects τ months lag or lead of treatment and cover all periods in the panel of length T . Event study models can only identify the full set of program effects $\{\beta_{\tau}\}_{\tau=-T}^T$ up to a constant. That is, event study models can only identify changes in program effects relative to a baseline level.¹⁶ In my preferred specification, I define the baseline as the second year before each household undertakes their initial conservation challenge. Using a baseline of the second year allows important pre-treatment trends in the 12 months preceding participation in the program to be estimated. Results are robust to other baseline periods; see Figure B.8 in the Appendix for a baseline of the third year pre-treatment. For the baseline of the second pre-treatment year, I estimate equation (1) while defining $\{\beta_{\tau} \equiv 0\}_{\tau=-12}^{-23}$. The non-parametric estimates $\hat{\beta}_{\tau}$ identify the average percentage change in monthly electricity

¹⁵For example, consider an observation in December 2009 for a household that began treatment (began its initial conservation challenge) in October 2009. For this household and observation, $t = \text{December}2009$ and $t - \tau + 1 = \text{October}2009$. This finds $D_{i,t-\tau+1} = 1$ only for $\tau = 3$.

¹⁶This can be seen by adding a constant to all β_{τ} and noting that this constant is then collinear with both the full set of individual and date fixed effects.

use within a household relative to their average electricity use during the baseline period.

Identification Challenges

Estimating causal effects in the event-study model requires treatment to be as good as randomly assigned conditional on time-invariant characteristics. This fails if participant and non-participant households would not have parallel trends in the absence of treatment. Evidence that the parallel trends assumption holds is obtained by comparing pre-treatment trends between participant and non-participant households. I estimate the following specification,

$$y_{it} = \beta_0 + \beta_1 m_t + \beta_2 m_t \times TPS_i + \beta_3 TPS_i + \epsilon_{it} \quad (2)$$

where y_{it} is log monthly electricity use, m_t is the date, and TPS_i is an indicator equal to one if household i is a Team Power Smart participant household. The only pre-treatment period available to all households is the year 2006. To test trends over multiple years I estimate specification (2) for several different time periods and include only participant households that do not begin a conservation challenge until 6 months after the initial pre-treatment years indicated. Table 2 shows the results where β_1 is the percent change per month for non-participant and participant households and β_2 is the additional monthly percent change for participant households. Participant households do not have a significantly different, at the 1% level, pre-treatment trend from non-participants. The magnitude of diverging time trends is also small. Taking the largest point estimate of different trends, $\hat{\beta}_2 = -0.00032$, would imply an upper bound to the potential bias in estimated program effects of only 0.4% at the end of the first conservation challenge.

Households starting a conservation challenge in response to past consumption shocks or expectations of future electricity use will also violate the parallel trends assumption. Shocks to past consumption can be tested by examining pre-treatment trends. As I discuss in Section 4.1, I find no evidence households begin a conservation challenge in response to increases in their past consumption such as large electricity bills from unusually cold winters. Whether households begin a challenge based on their expected future consumption cannot be tested. For example, it is not possible to distinguish between a household participating in a conservation challenge which causally reduces their electricity use from a household purchasing an efficient dryer and the latter participating to take advantage of their upcoming decline in electricity use. In this case, reductions in electricity use

Table 2: **Pre-Program Trends in Electricity Use**

	(1)	(2)	(3)	(4)	(5)
β_1 : Date	-0.00030** (0.00015)	-0.00034*** (0.00010)	-0.00034*** (0.00008)	-0.00061*** (0.00006)	-0.00053*** (0.00005)
β_2 : Date \times Participants	0.00032 (0.00021)	-0.00001 (0.00016)	-0.00003 (0.00013)	-0.00009 (0.00011)	0.00013 (0.00012)
Pre-Program Years	2	3	4	5	6
Non-Participant IDs	8877	8877	8877	8877	8877
Participant IDs	6778	5477	4609	3638	2162
Observations	375719	516743	647327	750898	794805

Notes: β_1 : *Date* is the pre-program time trend common to participant and non-participant households. β_2 : *Date \times Participants* is the additional time trend specific to participant households. Pre-Program Years is the length of time for which time trends are estimated and excludes all households that start a challenge within six months after the given pre-program. The six month period is to avoid pre-treatment trends that could include anticipation effects in the final months pre-treatment. Standard errors in parentheses clustered at the household level. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

that coincide with the start of participation in Team Power Smart would not be causal. Lacking a suitable instrument for initial participation or a discontinuity in eligibility for program participation such as used by Ito (2015), self-selection based on expectations of future electricity use cannot be ruled out. Instead, as I discuss below, the pattern of estimated program effects provides evidence that such self-selection is unlikely to be substantial and that the estimated reductions are primarily causally due to the Team Power Smart program.

4.1 Event Study Estimates—Initial Conservation Challenge

This section presents the results of estimating the event study model, equation (3), for all participant households and a comparison group of non-participant households.¹⁷ To visualize the trend of program effects over event-time τ , I plot the estimates $\hat{\beta}_\tau$ in Figure 1(a).¹⁸ The pattern of monthly estimates provides substantial insight into the short run and long run reductions in electricity use and whether reductions are causally due to the

¹⁷Due to the small sample size for months far from the initial conservation challenge I pool monthly indicators before and after the ± 5 year estimation window around the initial conservation challenge into separate indicators. Results are robust to this simplification. To ensure a closer comparison group of non-participants to participant households, in the results below I also use a random sub-sample of non-participant households with the same composition of building type and heating characteristics as participant households. Results are robust to including all non-participant households; see Figure B.7 in the Appendix.

¹⁸ Table E.8 of the Appendix lists all point estimates.

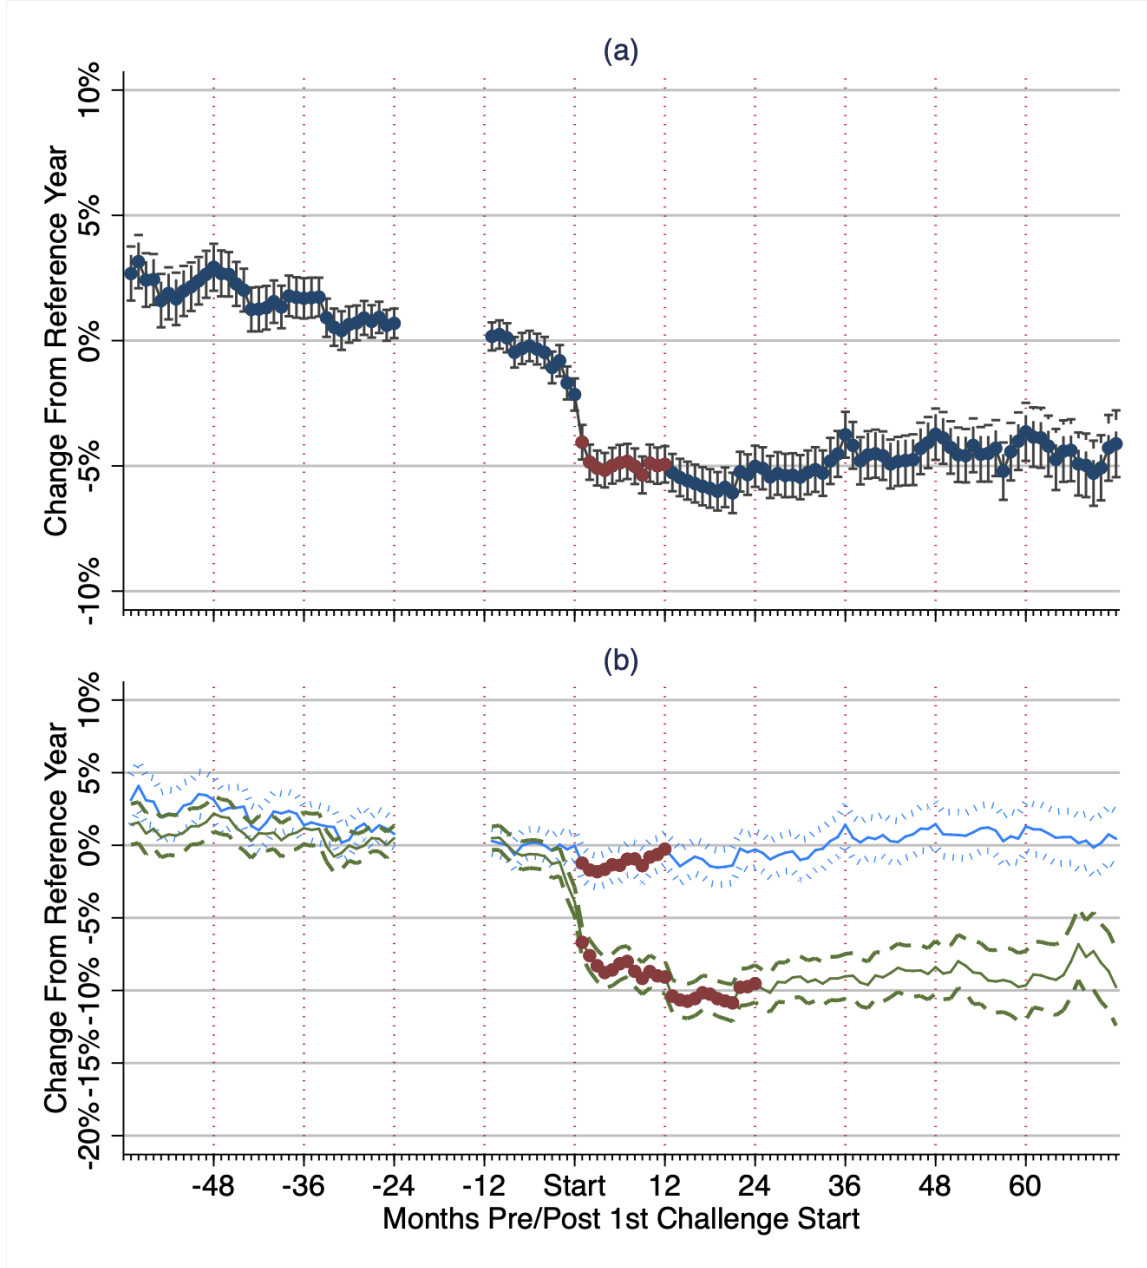
Team Power Smart program. The estimates in Figure 1(a) can be separated into three time intervals: the pre-treatment period leading up to the initial conservation challenge (months -60 to 0), the twelve months of the initial conservation challenge (months 1 to 12), and the months after the initial conservation challenge is completed (months 13 to 72).

In the second time interval, the twelve months of the initial conservation challenge show a substantial average reduction in electricity use of 4.9% relative to the baseline of the second year pre-treatment. BC Hydro measures households' reductions relative to the twelve months preceding a challenge; households achieve a 4.4% reduction relative to this year and their 10% conservation target. The lack of trend in monthly estimates of the initial challenge shows that, on average, households are not strongly increasing or decreasing their electricity conservation during the challenge. This stable conservation is consistent with households making either an initial investment that causes persistent reductions throughout the challenge, or maintaining a constant conservation effort throughout the challenge. The consistent program effects also indicate that self-selection to take advantage of upcoming short-term changes in electricity use, such as a holiday, is not a significant cause of participation; if they were, reductions in energy use would be expected to spike in the initial months of the challenge before partially rebounding.

In the first time interval, the months leading up to the initial conservation challenge show a gradual decline in estimated program effects. This trend could arise for two reasons. First, the pre-treatment trend could indicate a violation of the parallel trends assumption—differences in exogenous trends between the treated and control households—despite the lack of statistically significant different pre-trends estimated previously. To test this, I estimate specification (1) without non-participant control households. This identifies program effects by comparing currently treated households to a control group composed of households treated at a later date in the panel (Borusyak and Jaravel, 2016). Estimates lose precision due to the smaller set of households but continue to show the same pre-treatment trend, showing this decline is not due to a violation of the parallel trends assumption.¹⁹ Second, the pre-treatment decline could reflect different time trends among participant households that are not fully controlled for by common date fixed effects. Figures 1 and 2 plot the program effects for households undertaking different numbers of challenges. This shows that the declining pre-treatment trend is limited to households undertaking only one or two challenges and, importantly, does not bias the

¹⁹I plot estimates in Appendix Figure B.6.

Figure 1: **Estimated Program Effects:** All households & re-enrolling households



Notes: Estimated program effects $\hat{\beta}_\tau$ with 95% confidence intervals from equation (3). The pre-treatment period is denoted by the months prior to *Start* ($\tau \leq 0$). The visual gap in estimates between months $\tau = -11$ and $\tau = -23$ is the excluded reference period. $\hat{\beta}_\tau$ identify the percent change in electricity use relative to the average electricity use within a household during this excluded reference year. Estimates include individual and date fixed effects and I cluster standard errors at the household level. Points estimates are in the Online Appendix. (a) Estimates, including the period after the initial challenge concludes ($\tau > 12$), contain both households ending participation and those that re-enroll. (b) Estimates, plotted here as a line for clarity, in blue are households that undertake a single challenge and then end their program participation. Estimates in green are for households that undertake at least two conservation challenges and continue to their second challenge within 12 months of completing their initial challenge. Not shown are estimates θ_g for electricity use during the gap between the first and second challenges. Months 13-24 are estimates of the average change in electricity use among households in their second conservation challenge independent of any gap between challenges.

short or long-run program effect estimates.

A threat to estimated program effects not being causal is households beginning a conservation challenge in response to a high electricity bill, such as after a cold winter. In this case, reversion to the mean would result in reductions in electricity use, relative to the previous year, being credited to the program. If this self-selection occurs it will manifest itself as positive pre-treatment effects in the months immediately prior to the initial conservation challenge. Instead, the opposite effect is found; the last two months (months $\tau = [-1..0]$) in Figure 1 show a decline in electricity use relative to the previous months. These reductions in electricity use before the program starts do not count towards the conservation credited to households and, in addition, make it harder to achieve the financial reward as they lower the baseline from which the 10% conservation target is measured. This lack of a positive pre-treatment effect in the months prior to the challenge suggests that households do not self-select into the program based on past consumption shocks.

The pre-program decline of months $\tau = [-1, 0]$ may indicate that households are undertaking electricity conservation prior to the program start. This could indicate self-selection into Team Power Smart as a result of making an energy efficiency investment which would bias the estimated program conservation upwards. There are two reasons such self-selection and bias to the estimated program effects is unlikely to be large. First, participation in Team Power Smart as a result of making an investment would result in persistent reductions in electricity use; in contrast, subsection 4.2 shows that electricity use rebounds as households leave the program. Second, the pre-program roll-off can result mechanically from the billing process. During the period of the panel studied in this paper BC Hydro does not record electricity bills on a fixed monthly basis.²⁰ Instead, BC Hydro uses a rolling billing period where different houses are billed on different days of not more than 62 days and the use is calendarized to monthly consumption. As a result, reductions that occur after the start of a conservation challenge cannot be separated within a billing cycle from electricity use that occurred prior to the challenge start. This can result in reductions due to a conservation challenge being partially credited to up to the last two months before a household begins its challenge.²¹

²⁰BC Hydro began a rolling installation of Smart meters in 2011 which allowed collection of hourly electricity use. The event study estimates of this section includes households beginning their initial challenge prior to February 2013 and so will include some households with electricity use recorded at a higher frequency. However, all data provided has been harmonized and calendarized by BC Hydro to a monthly level regardless of the recording frequency. As a result, challenges over the period studied are recorded as beginning on the first day of a month regardless of when during the month a household signed up for a challenge, or when the electricity meter was read.

²¹Households could also potentially begin reducing their electricity use in response to a conservation

The third time interval includes all months after the initial challenge completes, months $\tau = [13..72]$. Pooling all participant households in estimating (1) includes those ending their program participation after their first challenge, those immediately continuing to additional challenges, and those waiting several months to years before starting a subsequent conservation challenge. Estimates $\hat{\beta}_\tau$, $\tau > 12$ are the average change in electricity use across these households, including any rebound in electricity use, and additional treatment effects from subsequent challenges undertaken by households. While Figure 1 shows that participation in Team Power Smart is associated with long-run average reductions in electricity use, this does not distinguish persistent energy savings due to a challenge, from the effects of subsequent challenges. To distinguish these effects I modify the event-study model (1) and estimate it separately for households ending their participation and those re-enrolling in subsequent conservation challenges.

4.2 Event Study Estimates—Additional Conservation Challenges

Comparing estimated program effects across households that undertake different numbers of conservation challenges sheds light on both the process of self-selection into additional challenges and the persistence of energy savings. To do this, I pool households by the number of conservation challenges they undertake and estimate separate event study models for each group of households, equation (3)

$$y_{it} = \alpha_i + d_t + \sum_{\tau=-59}^{72} \beta_\tau D_{i,t-\tau+1} + \sum_{g=1}^8 \theta_g G_{itg} + Pre_{it} + Post_{it} + \epsilon_{it} \quad (3)$$

where y_{it} , α_i , d_t , and $D_{i,t-\tau+1}$ are defined as in the event-study model of equation (1). I pool monthly indicators before and after the ± 5 year estimation window around the initial conservation challenge into, respectively, indicators Pre_{it} and $Post_{it}$. Some households undertaking multiple conservation challenges have a gap in time between when their previous challenge completes and they re-enroll in a subsequent challenge.²² To account for this variable gap I include the indicator G_{itg} in equation (3). G_{itg} is 1 if household i in month t has completed challenge g but has not yet re-enrolled in challenge $g + 1$. This indicator is not necessary for estimating the event study model; instead, it simplifies the comparison

challenge but not complete their online registration until some weeks later.

²²Appendix Figure A.3 shows histograms of these gaps.

of program effects across households with different gap lengths. For households undertaking two or more challenges, including G_{itg} defines the coefficients $\beta_{\tau, \tau = [12..23]}$ as the program effects of the twelve months of the second conservation challenge regardless of the length of gap between conservation challenges. For households ending their participation after a single challenge, $\beta_{\tau, \tau = [12..23]}$ are the post-program effects for the first 12 months immediately following the initial challenge. Similarly, including G_{itg} defines $\beta_{\tau, \tau = [24..35]}$ as the program effects during a third conservation challenge for households that re-enroll, and as the second year post-program effects for households that end their participation. In my preferred specifications I pool all households that re-enroll within 12 months of completing their previous challenge and exclude households with longer gaps between challenges; results are robust to alternate restrictions on challenge gap lengths.

Estimates for Additional Conservation Challenges

I separately estimate equation (3) for six subsets of households depending on how many conservation challenges they undertake; these subsets are not mutually exclusive.²³ In Figure 1(b) I plot estimates $\hat{\beta}_{\tau}$ for households that end their program participation at a single challenge, and separately for households that re-enroll in a second conservation challenge. These two household groups show significant differences. Households that end their program participation after the initial challenge have average reductions during the challenge of 1.2%. Households that re-enroll in a second conservation challenge have average reductions during their first challenge of 7.2%. This suggests self-selection into additional challenges based on the reductions in electricity use achieved during the initial conservation challenge. Alternatively, it could be that households which are ex-ante likely to continue to additional conservation challenges are also those households that achieve large reductions in energy use.

Estimates $\hat{\beta}_{\tau}$ in Figure 1(b) for households that re-enroll are the average program effect including both households ending their participation after their second conservation challenge, and those that re-enroll in a third challenge. Figure 2(a) separates these households and plots estimates $\hat{\beta}_{\tau}$ for those undertaking only two conservation challenges against those that re-enroll in a third challenge. Both these household groups have similar conservation during their first challenge. However, during their second conservation challenge those households that end their participation show a rebound in their use over the months of the challenge. Their electricity use returns close to the pre-program levels as the chal-

²³Point estimates are presented in Table E.9 in the Appendix.

challenge ends. In comparison, households that continue to a third conservation challenge continue to decrease their electricity use during their second challenge.

Figure 2(b) shows this pattern repeats again between households ending program participation after the third challenge and those re-enrolling in a fourth challenge.²⁴ Electricity use remains similar across both groups of households until the last challenge, it continues to decrease among those who re-enroll and rebounds among those who leave the program. This rebound does not return to the pre-program use; electricity use remains persistently lower than pre-program levels by approximately 6.5%.²⁵ In addition, the rebound occurs during the months leading up to the end of their final conservation challenge instead of after the challenge has completed. This is consistent with households stopping their conservation effort, or “giving up,” prior to the end of the challenge.

4.3 Program Effects by Household Characteristics and Season


Electricity conservation may vary by season depending on how households reduce their electricity use. Improved insulation, smart thermostats, or reductions in the household temperature will produce larger energy savings in the winter and among households that heat primarily with electricity. More efficient dryers, lightbulbs, or other changes that affect primarily non-seasonal electricity use will generate energy savings year round for both household heating types. Comparing when during the year that reductions in electricity use occur, and between household heating types, sheds light on how households have responded to the conservation challenge.

I estimate the event study model (1) separately for each calendar month and household heating type, and use the year immediately before the initial conservation challenge as the reference baseline. In Figure 3, I plot the estimated program effects for each month of year during the initial challenge. Both household heating types have similar reductions in the summer months, while Electric Space Heating households have larger reductions in the winter months. To estimate the fraction of reductions due to heating, I use the reductions in electricity use over the four warmest summer months as a measure of changes in non-heating use. Comparing this to the reductions over the remainder of the year

²⁴To clarify the changes, Appendix Figure B.9 plots estimates from Figures 1 and 2 together in the same graph.

²⁵This does not necessarily imply treatment effects were partially persistent. Idiosyncratic but persistent shocks to electricity use, combined with self-selection based on program success, can generate long-run lower electricity use among those that repeatedly re-enroll.

Figure 2: **Estimated Program Effects: Two, Three, and Four Challenges**



Notes: Estimated program effects $\hat{\beta}_\tau$ and 95% confidence intervals from equation (3) estimated for mutually exclusive groups of households. Estimates in blue are households that end their participation and estimates in green are households that re-enroll in an additional challenge. Estimation sample restricted to households that continue to subsequent challenges within 12 months. Not shown are estimates θ_g for electricity use during the gap between challenges. Estimates include individual and date fixed effects and I cluster standard errors at the household level. (a) compares households at the end of the second challenge, and (b) at the end of the third challenge.

finds that 18% of electricity conservation among non-electric heating households is related to heating, while 52% of the conservation among electric heating households is due to changes that affect heating. That non-electric heating households have reductions due to heating is not unexpected; non-electric heating households may still use after-market baseboard heaters in addition to their non-electric primary heat source. Figure 3 also plots estimated program effects for Electric and non-Electric heating households, showing they have similar conservation trends.

To estimate the program effects over the initial challenge across household characteristics I use an event study model with annual event-time indicators,


$$y_{it} = \sum_{\mathcal{T}=-4..-1,1..5} \theta_{\mathcal{T}} D_{i,t,\mathcal{T}} + \alpha_i + d_t + Pre_{it} + Post_{it} + \epsilon_{it} \quad (4)$$

where y_{it} is log monthly electricity use for household i at monthly date t , and $D_{i,t,\mathcal{T}}$ is an indicator for if household i in monthly date t is in year \mathcal{T} pre or post the challenge start date. α_i and d_t are household and monthly date indicators. Pre_{it} and $Post_{it}$ are indicators for households outside the ± 5 year window. θ_1 is the average change during the initial conservation challenge, relative to the average during the reference year. I use the last pre-program year as the reference year, $\mathcal{T} = 0$.

as well as pre-determined electricity use. To estimate the average change during a conservation challenge,

In Table 3 Panel A, I find that treatment effects in percentage terms are not statistically different across quartiles of pre-program use. This shows that reductions in absolute electricity use are larger for high-consumption household, but are not statistically different in percentage terms at a 5% level. In panel B, I find Apartments have larger, and statistically significant at a 5% level, reductions in electricity use compared to other dwelling types. In panel C I find no statistically significant differences, at the 5% level, in reductions by quartiles of assessed property values. If higher assessed value is taken as a proxy for income, the results imply that households respond similarly across quartiles of wealth. Panel D shows that dwellings in the top two quartiles of floor space have smaller reductions, in percentage terms, than smaller dwellings.

Figure 3: Program Effects By Heating Type



Notes: *Non-Electric Heating* are households expected by BC Hydro to heat primarily from sources other than electricity. *Electric Space Heating* are households whose primary source of heating is expected to be from electricity. The top panel plots estimated changes in electricity use during the initial conservation challenge relative to the same month in the year prior. Estimates are from equation (1) estimated separately for each month of the year. 95% confidence intervals shown by the dashed lines. The bottom pane plots estimates of $\hat{\beta}_\tau$ and 95% confidence intervals from specification (3) for all participant and non-participant households by heating type.

Table 3: Program Effects by Pre-Determined Variables

Panel A: Quartiles of Pre-Program Electricity Use				
	1st	2nd	3rd	4th
θ_1 : <i>Initial Challenge</i>	-0.0612*** (0.00658)	-0.0559*** (0.00533)	-0.0497*** (0.00492)	-0.0604*** (0.00482)
Avg. Use in 2006 (kWh)	370	648	917	1520
# of participants	1254	1253	1253	1253
# of non-participants	2220	2219	2219	2219
Panel B: Building Type				
	1 Sty SFD	2 Sty SFD	Apartment	Townhouse
θ_1 : <i>Initial Challenge</i>	-0.0432*** (0.00306)	-0.0381*** (0.00350)	-0.0625*** (0.00676)	-0.0396*** (0.00509)
# of participants	3426	2422	1282	1102
# of non-participants	3426	2422	1282	1102
Panel C: Quartiles of Assessed Value				
	1st	2nd	3rd	4th
θ_1 : <i>Initial Challenge</i>	-0.0465*** (0.00421)	-0.0474*** (0.00373)	-0.0409*** (0.00381)	-0.0396*** (0.00405)
Avg. Assessed Value (\$1,000)	\$288	\$479	\$686	\$1,210
# of participants	2220	2222	2217	2218
# of non-participants	2221	2218	2219	2219
Panel D: Quartiles of Floor Area				
	1st	2nd	3rd	4th
θ_1 : <i>Initial Challenge</i>	-0.0538*** (0.00467)	-0.0517*** (0.00383)	-0.0349*** (0.00362)	-0.0340*** (0.00403)
Avg. Floor Area (sq. ft.)	984	1639	2196	3275
# of participants	2129	2133	2129	2124
# of non-participants	2145	2129	2137	2134

Notes: Estimated average change in electricity use from the year pre-program to year of the initial conservation challenge. Panel A: Quartiles of pre-program electricity use determined from households' average electricity use in the pre-program year, 2006. Quartiles are defined separately for the balanced set of participant and non-participant households. Estimates exclude households starting their initial challenge before 2009 to avoid biasing estimates with a reversion to the mean. Panel B: Building type includes the four principal housing types of single story single family dwellings, two story single family dwellings, apartments, and town homes. Panel C: quartiles of assessed value are from the 2010 BC Assessment for individual units and include both structure and land value. Panel D: quartiles of floor area are in square feet. Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

5 Re-enrollment In a Second Conservation Challenge

Section 4 documented how electricity use diverges between those who leave Team Power Smart and those that re-enroll in an additional conservation challenge. This suggests that energy use may rebound in the absence of program participation and that voluntary re-enrollment in additional conservation challenges could cause additional reductions in electricity use. This section uses the same Team Power Smart program to study two aspects of these households’ extensive margin re-enrollment decisions; the decision to re-enroll itself, and causal reductions in electricity use from continuing to a second conservation challenge.

Weather affects households’ electricity use. BC Hydro applies a weather-adjustment algorithm to avoid penalizing or unnecessarily rewarding households for changes in electricity use that are due to idiosyncratic changes in weather, and not their conservation effort.²⁶ This algorithm is applied to the changes in actual electricity use that customers are billed for; I refer to these changes, prior to their weather-adjustment, as billed changes. The weather-adjustment algorithm adjusts billed changes for year-to-year changes in heating degree days, and results in a second measure which is the electricity conservation houses receive credit for; I refer to these as credited changes. When households view their online progress towards their 10% reduction goal, or the reductions in electricity use they are credited with achieving during a challenge, they are shown the credited — not billed — changes.

Figure 4 plots the probability of re-enrolling in a second challenge against the credited changes in electricity use from a household’s initial conservation challenge. This provides insights into households’ decisions whether to re-enroll. Among households that fail their initial challenge, their probability of continuing to a second challenge is largely independent of their credited changes in electricity use. Figure 4 shows households with large increases in electricity use of around 20% have a similar probability of re-enrolling as households that had no change, and these households are approximately 7% less likely to re-enroll as those that nearly achieved their 10% target. A similar pattern repeats amongst households that passed their challenge; those that barely pass with reductions $\sim 10\%$ are equally likely to continue as households that achieved reductions of $\sim 20\%$ or more. In contrast, there is a sharp discontinuous jump in the probability of continuing to a subsequent challenge at the 9.5% threshold for success.²⁷ This pattern shows that in

²⁶I discuss the algorithm in Appendix A.4.

²⁷While similar discontinuities exist after the second and subsequent conservation challenges, the sample

deciding whether to continue in the program, households are responsive to their success or failure in a conservation challenge but are largely insensitive to the level of reductions in electricity use they are credited with or achieve.²⁸

This insensitivity of re-enrollment to changes in electricity use is not consistent with the incentive structure of the program. Each conservation challenge is an additional 10% reduction compared to the year prior; this creates a greater incentive for households that had increases, or smaller reductions, in electricity use to re-enroll compared to if they had larger reductions. Households with large increases during a challenge establish a new higher baseline for their subsequent 10% conservation challenge, while those that decreased their electricity use establish a new lower baseline. Under the reasonable assumption that the marginal cost of electricity conservation is increasing, the greater the electricity conservation achieved in a challenge the more costly, in terms of effort or financial investment, the subsequent challenge will be.^{29,30}

The magnitude of the discontinuity in Figure 4 is large compared to differences in re-enrollment probability across household characteristics. To explore what correlates with households' re-enrollment decisions I estimate several Probit models—specification details are in Appendix C and results in Appendix Table C.1. These show that a households' pre-determined characteristics have little direct, or in-direct through level of electricity conservation, correlation with the probability of re-enrolling. Taking the largest difference in point estimates finds Townhouses are 8.6% more likely to re-enroll than homes classified Other, and households with Pre-Program electricity use three standard deviations above the mean 2006 use are 5.9% more likely to re-enroll. In comparison, households that pass their conservation Challenge are 20.2% more likely to re-enroll.


is too small for instrumental variable estimates.

²⁸See Figure C.1 in the Appendix and Table 4 for the comparison to billed reductions.

²⁹This is supported by Figure C.1 in the Appendix, which shows the fraction of households that, conditional on re-enrolling, pass their second conservation challenge. Consistent with increasing marginal cost of electricity conservation described above, approximately half of households with billed increases in electricity use of 10-20% pass their second challenge; these households are twice as likely to pass as households that had reductions during the first challenge of 10% or larger. This finding also holds when using credited changes.

³⁰A potential concern is that households may differ substantially depending on their credited changes in a way that affects their probability of re-enrolling, and this could offset the incentive to re-enroll caused by smaller reductions. The large weather adjustment makes such an offset unlikely to produce the observed insensitivity of re-enrollment to conservation. The weather adjustment significantly randomizes households achieving the same actual reduction in electricity use across the credited changes, increasing the comparability of households with different credited changes.

Figure 4: Probability of Re-Enrolling



Notes: Credited changes are the annual changes in electricity consumption between the pre-program baseline year and the year of the first conservation challenge after processing by the BC Hydro weather adjustment algorithm. Credited changes are those changes shown to households during their conservation challenges and after their completion. Point estimates and 95% confidence intervals are the mean probability of re-enrolling among households within 0.75% width binds in credited changes. The dashed line is a first order local polynomial fit. The vertical dashed line indicates the 9.5% threshold. By definition, households to the left of the dashed line pass their conservation challenge while those to the right fail.

5.1 Fuzzy Regression Discontinuity Empirical Strategy

I use the discontinuity in the probability of continuing to a second challenge at the 9.5% conservation threshold for success, shown in Figure 4, as the instrument for treatment in a second conservation challenge. The instrumental variable is a binary indicator for success in the initial conservation challenge. The first stage relationship is

$$C_i = \gamma_0 + \gamma_1 1\{R_i \leq \bar{R}\} + \gamma_2 R_i + \gamma_3 1\{R_i \leq \bar{R}\} \times R_i + \gamma_4 B_i + \gamma_5 X_i + \eta_i \quad (5)$$

where C_i is a binary indicator for whether a household continues to a second challenge, R_i are households' credited changes in electricity use from the first challenge, \bar{R} is the threshold for success in the challenge and is -9.5%, $1\{R_i \leq \bar{R}\}$ is the dummy variable for success in the initial challenge, B_i are the billed changes from the initial challenge, and X_i is a vector of other controls. In my main specification I control for a linear trend in credited reductions and allow this trend to have different slopes on either side of the discontinuity. The instrument excluded from the second stage is $1\{R_i \leq \bar{R}\}$.

The second-stage relationship is

$$y_i = \beta_0 + \beta_1 C_i + \beta_2 R_i + \beta_3 1\{R_i \leq \bar{R}\} \times R_i + \beta_4 B_i + \beta_5 X_i + \epsilon_i \quad (6)$$

where y_i is the post-challenge percent change in electricity use. y_i is defined

$$y_i \equiv \frac{(u_{i,\tau=2} - u_{i,\tau=1})}{u_{i,\tau=0}} \quad (7)$$

where $u_{i,\tau}$ is household i 's aggregate electricity use during the year indexed by event-time τ . For households that do not undertake a second challenge $u_{i,\tau=2}$ is the total electricity use in the 12 months immediately following the completion of their initial challenge, $u_{i,\tau=1}$ is the total electricity use during their initial challenge, and $u_{i,\tau=0}$ is the use during the pre-challenge year. For households that immediately undertake a second conservation challenge with no gap between challenges $u_{i,\tau=2}$ is the total electricity use during the second challenge and $u_{i,\tau=1}$ and $u_{i,\tau=0}$ are as before. For households that wait before beginning a second conservation challenge I define $u_{i,\tau=2}$ as the 12 months of electricity use during their second challenge and $u_{i,\tau=1}$ as the 12 months of electricity use immediately preceding that second challenge. This makes y_i a consistent measure of the reductions in electricity use a household is trying to achieve in its second challenge regardless of

whether that household waited before undertaking a challenge or began it immediately. I center the billed and credited changes at the 9.5% threshold. β_0 is the post-challenge change in billed electricity use at this threshold for households that do not continue in the program. β_1 is the additional effect on post-challenge billed changes in electricity use relative to households that left the program.

5.2 Fuzzy Regression Discontinuity Identifying Assumptions

A regression discontinuity strategy requires that units on one side of the threshold that defines treatment are a suitable counterfactual for units on the other. (Lee and Lemieux, 2010) This assumption would be violated if households can precisely manipulate their assignment into treatment, and such manipulation results in households sorting at the discontinuity in any way that affects potential outcomes. Such sorting is a concern in this setup as households are explicitly trying to achieve a 10% conservation target. Sorting could occur if, for example, households are heterogeneous in their attention to their progress and high-information type households exert additional effort in the last months of a conservation challenge and self-select into passing their challenge.

Fortunately, the separation of the 9.5% threshold for Challenge success from the 10% reduction target allows me to show that this RD strategy remains valid despite households manipulating the running variable of electricity conservation such that they exhibit bunching below the 10% reduction target. Sorting discontinuously at the 9.5% threshold for success is unlikely for several reasons. Most importantly, households were not aware that their success or failure would be evaluated against a 9.5% threshold instead of the advertised 10% target. This does not remove the potential problem of sorting; households sorting around the 10% target in a way that changed potential outcomes in the absence of treatment would invalidate the causal interpretation of regression discontinuity estimates at the 9.5% threshold.³¹ In addition, the weather adjustment mechanically randomizes households near the threshold into and out of treatment based on the weather change that occurred. The cumulative nature of the challenge also makes precise manipulation of success difficult. A household at a 9% cumulative reduction entering the last month of their challenge would have to double their previous monthly reductions and reduce their use in the last month by 21% to achieve their 10% target.

³¹In theory, a sufficiently large number of observations would allow estimation limited to bandwidths of $\pm 0.5\%$ around the 9.5% threshold, thus avoiding the problem of sorting around the 10% threshold but changing the subset of households that the LATE estimates are estimated for.

Evidence on whether there is sorting on observables can be gained by testing the continuity of covariates and pre-determined variables across the 9.5% threshold and 10% target. I find no statistically significant change which supports that households are not sorting at the 9.5% discontinuity or 10% threshold in a manner correlated with observables.³² This is supported by a density test of McCrary (2008), which fails to reject the null hypothesis (one sided p-value 0.117) of no discontinuity in the density of the running variable at the 9.5% threshold.³³

However, a McCrary (2008) density test rejects that there is no discontinuity in the density of electricity conservation at the 10% threshold—suggesting that households can manipulate their assignment around the 10% target.³⁴ Such manipulation only invalidates the RD strategy if causes discontinuous potential outcomes. While the continuity of covariates demonstrated above suggests sorting on observables is not occurring, BC Hydro’s separation of the 10% target from the 9.5% threshold allows the continuity of potential outcomes at the 10% target to be directly tested. If households were sorting around the 10% target in a manner that affected potential outcomes, their decision to re-enroll in a second challenge or their post-program electricity use would be expected to be discontinuous. This does not occur—see C.1 for details. As a result, the discontinuity in Figure 4 is due to households succeeding in their conservation challenge and is not due to manipulation of their electricity conservation that results in sorting around the 9.5% threshold.

The exclusion restriction requires that the instrument only affect the outcome, y_i , through the decision to continue to a second challenge (Angrist and Pischke, 2008). Conditional on credited and billed reductions, success in a challenge can have no direct effect on post-program changes in electricity use and therefore can be excluded from the second stage. This assumption could be violated if households receive a warm-glow effect from succeeding that affects their subsequent effort at reductions independent of continuing, or if the \$75 rebate causes an income effect and alters post-program conservation. As \$75 is small relative to household’s incomes I assume there is no income effect that influences electricity use. Any warm-glow effect on subsequent conservation effort is likely to be short lived compared to the effect of the financial incentive, which remains throughout

³²See Appendix Table C.1 and Appendix Figure C.1.

³³See Appendix Figure C.3.

³⁴Appendix Figure C.2 shows a histogram of credited changes during households’ first conservation challenges. The mass of observations just below the 10% threshold suggests that households may be bunching around the 10% target. This is tested in Appendix Figure C.2 which rejects the null of no discontinuity with a one sided p-value of 0.0017.

the twelve months of the challenge. In addition, if a warm glow effect was substantial it is likely to be particularly strong during the initial months of the next challenge while a household’s success is still fresh in their minds. The event study estimates of Figure 1(b) indicate that additional program effects during the second challenge are consistent throughout the twelve months of challenge, suggesting that there is little warm-glow effect.

5.3 Fuzzy Regression Discontinuity Estimates

This section presents the fuzzy-RD estimates of the treatment effect of a second conservation challenge. Across a wide variety of specifications and robustness checks I find a consistent pattern where re-enrolling in a second conservation challenge causes a large additional reduction in electricity use. These results are consistent with the event-study results and support that additional conservation challenges cause additional reductions in electricity use, and that electricity use rebounds as households leave the program.

Table 4 presents my preferred specification. Columns (2) through (6) show results estimated for different bandwidths from $\pm 7\%$ to $\pm 3\%$ around the threshold of a -9.5% change in credited electricity use.^{35,36} Panel (A) shows the first-stage results for the probability of continuing to a second challenge, estimated from equation (5). For my preferred bandwidth of 5% I find that, conditional on failing the challenge, 53% of households continue to a second conservation challenge. At the 9.5% threshold for success, households that just succeed in their initial conservation challenge are 14.5% more likely to continue to an additional challenge than those which just failed. This pattern repeats across various estimation window widths. Approximately 50% of households re-enroll in a second conservation challenge if they fail their initial challenge, while success in a challenge causes an additional 14% to 20% of households to re-enroll.³⁷ Across bandwidths from $\pm 7\%$ to $\pm 3\%$ I find the F-statistic on the instrument decreases from 22 to 6.5. This indicates that the first-stage is reasonably strong for larger bandwidths, but the small sample size becomes relevant as the bandwidth narrows.

Table 4 panel (B) reports the OLS and second-stage instrumental variable estimates of

³⁵I present estimates using a range of bandwidths and a uniform weighting instead of kernel estimates. See [Imbens and Lemieux \(2007\)](#) for a discussion on the practical similarities of varying the bandwidth to using different kernels. Results are robust across different bandwidths when using a triangular kernel.

³⁶Plots of First Stage and Reduced Form with 7% bandwidths are presented in Appendix Figure C.4. 5% and 9% bandwidths are presented in Appendix Figure C.5.

³⁷A weak instruments test by [Moreira \(2003\)](#) rejects (p-value 0.013, 5% bandwidth) that the binary indicator for Success (γ_1) is a weak instrument.

equation (6). Specification (1) is the OLS result for all households. (1) shows that re-enrolling in an additional challenge is associated with a 1.6% decline in post-challenge electricity use, relative to households that do not re-enroll. Specifications (2) through (6) show the IV estimates for different estimation bandwidths. These estimates find a consistent pattern where, for households that comply with the instrument, continuing to a second conservation challenge causes a reduction in electricity use. For my preferred bandwidth of 5%, I estimate that continuing to a second conservation challenge causes a 23% reduction in electricity use ($\hat{\beta}_1 = -0.231$). This is a large effect. By definition, the treatment effect is the change in electricity use for complier households relative to what they would have had, had they not re-enrolled. An instrumental variables estimation strategy cannot identify the level of electricity use, in the absence of continuing, for these compliers. As a result, the treatment effect does not separately identify a potential rebound in electricity use among those that leave the program from additional reductions in electricity use beyond those achieved in the first conservation challenge.³⁸ Given the average reduction of 9.1% during the first challenge for households within the $\pm 5\%$ bandwidth, and the target of an additional 10% conservation, the estimated treatment effect is likely comprised of both a significant rebound in electricity use among complier households that end participation and a large additional reduction among complier households that re-enroll in another challenge.³⁹

In [Appendix D](#) I undertake a variety of robustness checks. These include: 1st and 2nd order bias-corrected estimates and optimal bandwidths selected using the variance-bias tradeoff method of [Calonico et al. \(2014\)](#), restricting the sample to households experiencing small weather shocks, using additional covariates, using an alternate challenge gap length of 6 months, and an alternative specification using the log of monthly electricity use. While some estimates lose significance, especially for small bandwidths, in all cases the point estimates stay a consistent sign and large magnitude. This suggests that while the magnitude of point estimates varies, the causal effect of an additional conservation challenge is a large additional reduction in electricity use for those households whose decision to re-enroll in Team Power Smart is affected by their success or failure in their prior conservation challenge.

³⁸ $\hat{\beta}_0$ is not the average electricity use, conditional on other covariates, for complier households in the absence of a conservation challenge.

³⁹In comparing the OLS and fuzzy-RD estimates in Table 4 panel (B), it is important to note that the OLS estimates are for all households within the estimation window while the fuzzy-regression discontinuity design is a LATE for compliers. As a result, the difference in estimates may be partially due to compliers being responsive to the incentive as estimated, with always-taker's and never-taker's electricity conservation remaining largely unresponsive to the financial reward incentive.

Table 4: **Fuzzy Regression Discontinuity Estimates of a Second Challenge**

	(1)	(2)	(3)	(4)	(5)	(6)
Panel A – First Stage						
Dependent variable: Continue to a Second Challenge C_i						
Window		$\pm 7\%$	$\pm 6\%$	$\pm 5\%$	$\pm 4\%$	$\pm 3\%$
γ_1 : Success Ind.		0.202*** (0.0432)	0.190*** (0.0468)	0.145*** (0.0516)	0.137** (0.0573)	0.173** (0.0677)
γ_2 : Cred. Reduc.		-0.532 (0.778)	-1.391 (0.971)	-2.937** (1.244)	-1.290 (1.760)	2.090 (2.765)
γ_3 : Success \times Cred. Reduc.		1.166 (1.104)	2.491* (1.383)	2.565 (1.809)	-0.778 (2.447)	-3.973 (3.906)
γ_4 : Billed Reduc.		-0.300 (0.324)	-0.368 (0.344)	-0.0474 (0.366)	-0.282 (0.404)	-0.511 (0.491)
γ_0 : Constant		0.487*** (0.0303)	0.508*** (0.0330)	0.530*** (0.0364)	0.510*** (0.0409)	0.475*** (0.0479)
F-statistic		21.95	16.41	7.882	5.668	6.485
Panel B – Second Stage						
Dependent variable: Percent change in post-challenge electricity use						
Window	OLS	Instrumental Variable Estimates				
		$\pm 7\%$	$\pm 6\%$	$\pm 5\%$	$\pm 4\%$	$\pm 3\%$
β_1 : Re-Enroll	-0.0160*** (0.00422)	-0.125** (0.0605)	-0.178** (0.0738)	-0.231** (0.116)	-0.323** (0.164)	-0.183* (0.111)
β_2 : Cred. Reduc.		-0.412* (0.241)	-0.643* (0.354)	-1.185* (0.654)	-1.108 (0.828)	0.785 (0.654)
β_3 : Success \times Cred. Reduc.		0.303 (0.310)	0.375 (0.426)	0.867 (0.624)	-0.229 (0.985)	-1.732 (1.229)
β_4 : Billed Reduc.		-0.0461 (0.0917)	-0.113 (0.103)	-0.101 (0.121)	-0.216 (0.168)	-0.222 (0.161)
β_0 : Constant	-0.00773*** (0.00289)	0.0712** (0.0356)	0.103** (0.0447)	0.138** (0.0698)	0.184* (0.0950)	0.0944 (0.0625)
N	5432	2050	1763	1475	1196	888

Notes: This table reports fuzzy-RD estimates corresponding to equations (5) and (6). Estimation sample is restricted to households that either start their next challenge within 12 months or do not undertake an additional challenge. Estimation window is restricted to \pm the listed percent around the 9.5% threshold in credited changes. Standard errors in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

6 Cost Effectiveness

The Team Power Smart program is designed to produce electricity generation capacity savings and reduce the expected future increase in demand for electricity. A full welfare analysis of the TPS program is beyond the scope of this work. Instead, I provide a lower bound on the cost of avoided electricity generation caused by the program. I estimate a lower bound for two reasons. First, because the costs of administering and advertising the Team Power Smart reward program are confidential to BC Hydro, I consider only the cost of the \$75 rebates rewarded to households, and leave aside the costs of administering the program.⁴⁰ Second, I make the assumption that the estimated electricity conservation from the event-study model, Figure 1, are treatment effects; any overestimate of the true treatment effect will bias upward the cost of avoided generation. This is a reasonable assumption given the findings of Sections 4 and 5. From the estimated electricity conservation and the average electricity use among participants, Table A.1, I find that the average aggregate reduction in electricity use over the first six years after an initial challenge is 2.7 MWh per household. This is the average across all households and accounts for their decisions whether to re-enroll after each challenge. Taking into account households' average success in their conservation challenges, and the number continuing to additional challenges, the average aggregate rebate payment over the six years is \$53 per household. This finds an average cost of avoided generation of \$20/MWh. In comparison, participants paid an average retail price of \$96/MWh in 2016, while a large hydroelectric dam under construction in the province is estimated to have a levelized cost of electricity of \$34-\$83/MWh (British Columbia Utilities Commission, 2017). This makes the Team Power Smart program a cost-effective way to reduce the demand for electricity in comparison to the cost of new generation.

What is the cost of avoided carbon emissions due to this energy conservation program? It is important to note that this energy conservation program was not designed to principally reduce carbon emissions and the cost of avoided emissions is not particularly relevant to BC Hydro. British Columbia generates over 90% of its electricity from hydroelectric dams and has a low emissions intensity of 9kgCO₂e/MWh (BCH, 2016); \$20 per avoided

⁴⁰Because the program is administered online, variable costs excluding the rebate are likely to be negligible. Program fixed costs may not be insignificant relative to the cost of the rewarded rebates. One full time equivalent employee compensated at \$70,000 per year for managing the program would add approximately \$6 per challenge to the program. This is \$20 per awarded rebate using the 30% success rate over the initial five challenges households undertake.

MWh is a cost of avoided greenhouse gas emissions of \$2,222/tCO₂eq. However, BC Hydro engages in large cross-border trade in electricity with the United States, primarily California. Lower electricity use in British Columbia allows BC Hydro to sell relatively low cost and low emissions power to California. Assuming all reductions in B.C. electricity use reduces generation in California finds, using the 2017 California average emissions intensity (EPA, 2017), a cost of emissions abatement of \$71/tCO₂. At the 2017 U.S. average emission intensity, this falls to \$45/tCO₂ (Schivley et al., 2018)EPA (2016). These abatement costs are within the range of commonly discussed estimates of the SCC (EPA, 2016), and indicate that in some jurisdictions, repeated financial reward programs similar to the one studied in this work may be cost effective from a social perspective. The results of this paper show that the continued incentive of repeated financial rewards is important for maintaining and causing additional reductions in electricity use. The continued incentive improves the program’s cost-effectiveness, compared to a program offering a single annual conservation challenge. This improved cost-effectiveness occurs for two reasons. First, the program administration fixed costs are spread across additional conservation challenges. Second, the repeated incentive causes additional reductions, and keeps use from rebounding back close to pre-program levels.

7 Conclusion

This paper investigates how households respond to repeated financial rewards offered in an electricity conservation program. Using a panel of monthly electricity use, I track households’ decisions whether to re-enroll in the program along with changes in their electricity use. I use two separate empirical strategies to estimate changes in electricity use. These two strategies complement each other and allow me to address the potential for self-selection into the program, and into successive annual conservation challenges, that could bias estimated treatment effects. The ten years of the panel allow me to estimate the long-run persistence of electricity conservation as households leave the program, as well as the long-run changes among households that repeatedly re-enroll in the program. By comparing the electricity conservation across households I provide insights into what changes within the home households made to conserve electricity, and what information households use in deciding whether to re-enroll.

Two findings distinguish this work. First, this paper undertakes one of the longest studies of the persistence of electricity conservation. I do not find that financial rewards cause

persistent changes; instead, electricity use rebounds close to pre-program levels as households leave the program, and continues to decline among households that re-enroll. This shows that the ongoing incentive of successive financial rewards is necessary to cause long-run lower electricity use. A potential implication of this is that programs in fields outside energy conservation, such as education, that aim to cause persistent effects should consider incentivizing specific changes rather than reward a general outcome. Second, the repeated financial rewards provide an opportunity to observe both extensive and intensive margin decisions. I find that in deciding whether to re-enroll, households are responsive to their success or failure in achieving their conservation target, and yet, conditional on this success or failure, are notably unresponsive to their degree of electricity conservation. This suggests that households use heuristics in making participation decisions, rather than accounting for the incentive structure of the program or their degree of conservation effort.

This paper’s findings imply it is important for repeated incentive programs to explicitly consider self-selection into and out of re-enrollment. This is especially important when treatment effects fully dissipate after the incentive ends. Programs may be able to improve their cost effectiveness by encouraging re-enrollment, particularly when consumers fail to achieve their goal and self-select out of the program as a result. One way to address this could be to offer tiered incentives, such as offering an alternative challenge only to consumers that missed their target and are likely to self-select out of the program. Such an alternative incentive could be either transparently offered at the outset of participation, or revealed only to consumers that have not re-enrolled after some time period to target those likely to be ending participation. The responsiveness of households to their success or failure also opens up the possibility of modifying the reward structure to exploit this sensitivity; for example, by using unexpected consolation prizes instead of an all-or-nothing reward.

The use of information is central to how we model the decisions of agents across many questions within economics, from how tax rates affect labour market responses to voting behavior. Similar to households’ extensive margin decisions in this paper, the use of heuristics may be ubiquitous in intensive margin responses to price changes and many types of information. Heuristics could result from rational inattention and allow households to avoid the cognitive and time costs of considering complex incentives. Alternatively, they could result from households responding to incentives and information based on models of behaviour beyond standard neoclassical models of decisions. Disentangling the ways in

which consumers use information is an important avenue for future research. It is particularly important to the design of complex price schedules that are increasingly facilitated by the spread of smart meters, distributed electricity generation, and storage.

A further direction for inquiry are the consequences of repeating interventions. This paper showed how interventions that do not cause persistent effects can still produce cost-effective long-run changes by being repeatedly offered. These findings also suggest that consumers respond to each discrete reward in isolation, and do not consider the dynamic consequences of their effort on the future incentive structure. This raises the questions of whether this is a feature of the annual time scales of this program, or is a common feature of consumer responses to discrete rewards, and how consumers would respond as the time between potential rewards was reduced towards the limit of a continuous incentive.

References

- Allcott, H. (2011). Social norms and energy conservation. *Journal of Public Economics*, 95(9-10):1082–1095.
- Allcott, H. and Rogers, T. (2014). The Short-Run and Long-Run Effects of Behavioral Interventions : Experimental Evidence from Energy Conservation. *American Economic Review*, 104(10):3003–3037.
- Angrist, J. D. and Pischke, J. S. (2008). Mostly Harmless Econometrics : An Empiricist’s Companion.
- Bandura, A. (1977). Self-efficacy : Toward a Unifying Theory of Behavioral Change. *Psychological Review*, 84(2):191–215.
- BCH (2014). British Columbia Hydro and Power Authority Annual Report: 2014/2015.
- BCH (2016). BC Hydro Greenhouse Gas Intensities: 2007-2015. Available online at https://www.bchydro.com/about/sustainability/climate_action/greenhouse_gases.html (accessed July 15, 2018).
- Boomhower, J. and Davis, L. W. (2014). A credible approach for measuring inframarginal participation in energy efficiency programs. *Journal of Public Economics*, 113:67–79.
- Borusyak, K. and Jaravel, X. (2016). Revisiting Event Study Designs. Available online at <http://ssrn.com/abstract=2826228> (accessed July 15, 2018).
- British Columbia Utilities Commission (2017). British Columbia Utilities Commission Inquiry Respecting Site C: Preliminary Report to the Government of British Columbia. Technical report.

- Calonico, S., Cattaneo, M. D., and Titiunik, R. (2014). Robust data-driven inference in the regression-discontinuity design. *Stata Journal*, 14(4):909–946.
- Dolan, P. and Metcalfe, R. (2015). Neighbors, Knowledge, and Nuggets: Two Natural Field Experiments on the Role of Incentives on Energy Conservation. Available online at: <http://ssrn.com/abstract=2589269> (accessed July 15, 2018).
- EPA (2016). EPA Fact Sheet: Social Cost of Carbon. Technical Report December.
- EPA, C. (2017). California GHG Emission Inventory: California Greenhouse Gas Emissions for 2000 to 2015 - Trends of Emissions and Other Indictators. Technical report, Air Resources Board.
- Gerard, F. and Costa, F. J. M. (2015). Hysteresis and the Social Cost of Corrective Policies: Evidence from a Temporary Energy Saving Program. *Working Paper*.
- Gilbert, B. and Graff Zivin, J. S. (2013). Dynamic salience with intermittent billing: Evidence from smart electricity meters. *NBER Working Paper Series*, (Working Paper 19510).
- Gneezy, U., Meier, S., and Rey-Biel, P. (2011). When and Why Incentives (Don’t) Work to Modify Behavior. *Journal of Economic Perspectives*, 25(4):191–210.
- Imbens, G. W. and Lemieux, T. (2007). Regression discontinuity designs: A guide to practice. *Journal of Econometrics*, 142(2):615–635.
- Ito, K. (2014). Do consumers respond to marginal or average price? Evidence from nonlinear electricity pricing. *American Economic Review*, 104(2):537–563.
- Ito, K. (2015). Asymmetric Incentives in Subsidies: Evidence from a Large-Scale Electricity Rebate Program. *American Economic Journal: Economic Policy*, 7(3):209–237.
- Ito, K., Ida, T., and Tanaka, M. (2015). The Persistence of Moral Suasion and Economic Incentives: Field Experimental Evidence from Energy Demand. *NBER Working Paper Series*.
- Jessoe, K. and Rapson, D. (2014). American Economic Association Knowledge is (Less) Power: Experimental Evidence from Residential Energy Use. 104(4):1417–1438.
- Kassirer, J., Korteland, A., and Pedersen, M. (2014). Team Power Smart Sparks Increase in Low-Priority, Repetitive Behaviors. *Social Marketing Quarterly*, 20(3):165–185.
- Lee, D. S. and Lemieux, T. (2010). Regression Discontinuity Designs in Economics. *Journal of Economic Literature*, 48:281–355.
- Martin, S. and Rivers, N. (2016). Information provision, market incentives, and household electricity consumption: Evidence from a large-scale field trial. *Working paper*.
- McClelland, L. and Cook, S. W. (1980). Promoting Energy Conservation in Master

- Metered Apartments through Group Financial Incentives. *Journal of Applied Social Psychology*, 10(1):20–31.
- McCrary, J. (2008). Manipulation of the running variable in the regression discontinuity design: A density test. *Journal of Econometrics*, 142(2):698–714.
- Midden, C. J., Meter, J. F., Weenig, M. H., and Zieverink, H. J. (1983). Using feedback, reinforcement and information to reduce energy consumption in households: A field-experiment. *Journal of Economic Psychology*, 3(1):65–86.
- Mizobuchi, K. and Takeuchi, K. (2012). Using economic incentives to reduce electricity consumption: A field experiment in Matsuyama, Japan. *International Journal of Energy Economics and Policy*, 2(4):318–332.
- Moreira, M. J. (2003). A Conditional Likelihood Ratio Test for Structural Models. *Econometrica*, 71(4):1027–1048.
- Schivley, G., Azevedo, I., and Samaras, C. (2018). Assessing the evolution of power sector carbon intensity in the United States. *Environmental Research Letters*, 13(6).
- Schleich, J., Faure, C., and Klobasa, M. (2017). Persistence of the effects of providing feedback alongside smart metering devices on household electricity demand. *Energy Policy*, 107(April):225–233.
- Sexton, S. (2015). Automatic Bill Payment and Salience Effects: Evidence From Electricity Consumption. *Review of Economics and Statistics*, 97(2):229–241.
- Wichman, C. J. (2017). Information provision and consumer behavior: A natural experiment in billing frequency. *Journal of Public Economics*, 152:13–33.
- Winnett, R. A., Kagel, J. H., Battalio, R. C., and Winkler, R. C. (1978). Effects of monetary rebates, feedback, and information on residential electricity conservation. *Journal of Applied Psychology*, 63(1):73–80.

The Online Appendix is available at www.alastairfraser.ca